

Programmazione e Calcolo Scientifico

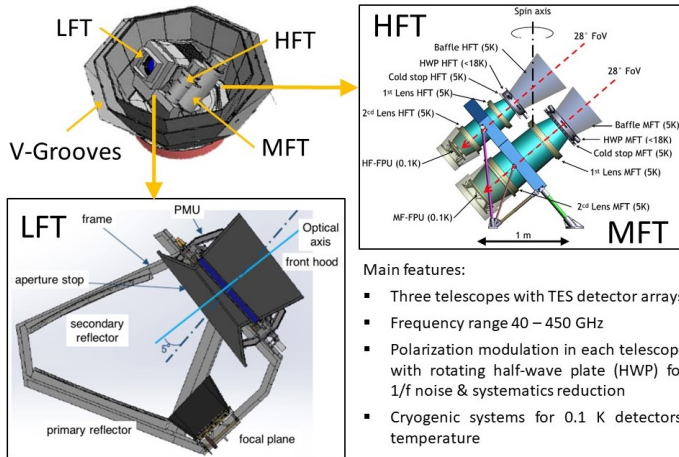
Matteo Cicuttin

Politecnico di Torino

March 9, 2024

What do we do with scientific computing?

LiteBIRD Payload Module



Main features:

- Three telescopes with TES detector arrays
- Frequency range 40 – 450 GHz
- Polarization modulation in each telescope with rotating half-wave plate (HWP) for 1/f noise & systematics reduction
- Cryogenic systems for 0.1 K detectors' temperature

What do we do with scientific computing?

LiteBIRD is a satellite that will be launched in 2029. It will be used to observe the CMB.


- Simulation of one of the LiteBIRD telescopes
- A couple of hours on 4 AMD MI250 GPUs

Click to play

LUMI supercomputer

www.top500.org/lists/top500/2023/06/

Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
1	Frontier - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE DOE/SC/Oak Ridge National Laboratory United States	8,699,904	1,194.00	1,679.82	22,703
2	Supercomputer Fugaku - Supercomputer Fugaku, AMD EPYC 9610 2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,630,848	442.01	537.21	29,899
3	LUMI - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE EuroHPC/CSC Finland	2,220,288	309.10	428.70	6,016



Course outline

In the era of Artificial Intelligence, we must be very well aware that computers are actually very stupid machines with **definite limits and capabilities**. **There's no magic**. **Santa Claus does not exist**.

We want to make you aware of those limits and give you some tools to maximally exploit computer capabilities.

Some topic we will cover:

- Anatomy of a computer (processor, memory, bus, cache)
- The limitations of a computer
- C++ programming
- Some simple algorithms and data structures

The AXPY operation

BLAS stands for Basic Linear Algebra Subprograms. It is a set of standard functions implementing the most common operations occurring in scientific computing:

- Level 1: vector-vector operations, $O(n)$ complexity
- Level 2: matrix-vector operations, $O(n^2)$ complexity
- Level 3: matrix-matrix operations, $O(n^3)$ complexity

The AXPY operation

BLAS stands for Basic Linear Algebra Subprograms. It is a set of standard functions implementing the most common operations occurring in scientific computing:

- Level 1: vector-vector operations, $O(n)$ complexity
- Level 2: matrix-vector operations, $O(n^2)$ complexity
- Level 3: matrix-matrix operations, $O(n^3)$ complexity

One of the operations included in BLAS is the AXPY: $\mathbf{y} \leftarrow \alpha \mathbf{x} + \mathbf{y}$ where $\mathbf{x}, \mathbf{y} \in \mathbb{R}^N, \alpha \in \mathbb{R}$.

```
for (size_t i = 0; i < N; i++)  
    y[i] = alpha*x[i] + y[i];
```

Anatomy of AXPY

Assume to run AXPY on a vector of size N:

```
for (size_t i = 0; i < N; i++)  
    y[i] = alpha*x[i] + y[i];
```

Observations:

- N cycles
- For each cycle: 1 multiplication and 1 addition

Total cost: $2N$ FLOPS (FLOP = Floating point Operation).

Let's measure how much time an AXPY takes and determine speed in FLOPS/s.

An hypothetical XNPY operation

We define XNPY as follows (XNPY is **NOT** part of BLAS). Let \mathbf{x}, \mathbf{y} be two vectors of equal length:

$$y_i \leftarrow x_i^n + y_i \quad i \in 1 \dots \text{length}(\mathbf{y})$$

Translated to code:

```
for (size_t i = 0; i < N; i++) {  
    double xpow = 1.0;  
    for (size_t p = 0; p < pow; p++)  
        xpow *= x[i];  
    y[i] = xpow + y[i];  
}
```

For each **outer** cycle we do $\text{pow}+1$ FLOPS.

Let's measure the speed of XNPY.

Any idea to improve AXPY?

Why AXPY is so slow compared to XNPY? Any idea to improve it?

```
for (size_t i = 0; i < N; i++)  
    y[i] = alpha*x[i] + y[i];
```

Matrix multiplication

Let $\mathbf{A}, \mathbf{B}, \mathbf{C} \in \mathbb{R}^{N \times N}$. The matrix multiplication $\mathbf{C} = \mathbf{AB}$ is computed as $c_{ij} = a_{ik} b_{kj}$. This is probably the most important operation in scientific computing.

The code for the matrix multiplication looks like this:

```
for (size_t i = 0; i < N; i++)  
    for (size_t j = 0; j < N; j++)  
        for (size_t k = 0; k < N; k++)  
            c[I(i,j)] += a[I(i,k)] * b[I(k,j)];
```

Total num of FLOPs is $2N^3$. Let's do some measurements!

Another matrix multiplication

We reorder the operations in our matrix multiplication:

```
for (size_t i = 0; i < N; i++)  
    for (size_t k = 0; k < N; k++)           // swapped this  
        for (size_t j = 0; j < N; j++)       // and this  
            c[I(i,j)] += a[I(i,k)] * b[I(k,j)];
```

Total num of FLOPs is **again** $2N^3$. Let's do some measurements!

What happens to performance?

- Why the performance of XNPY is better than AXPY despite being more complex?
- Why parallel XNPY improves but parallel AXPY does not?
- Why swapping two cycles in matrix multiplication changes the execution time so much?
- More generally, how do you choose a computer to solve a specific problem?

What happens to performance?

- Why the performance of XNPY is better than AXPY despite being more complex?
- Why parallel XNPY improves but parallel AXPY does not?
- Why swapping two cycles in matrix multiplication changes the execution time so much?
- More generally, how do you choose a computer to solve a specific problem?

The answer to these questions is rooted deeply

- in the way **computers** work
- in the way **compilers** work

Some tales about floating point

Before being **fast**, code must be **correct**.

Your code can go infinitely fast if you don't care about correctness, but people won't be interested in offering you a good job.

I - The Chaotic Bank Society

Your bank proposes to you to open a new account with the following scheme:

- When you open the account, you deposit $e - 1$ euros
- Each year the bank multiplies your savings by the number of years, but it takes a fee of 1 euro.

$$\begin{cases} b_0 = e - 1 \\ b_i = i \cdot b_{i-1} - 1 \end{cases}$$

You make a little program to figure out what happens after 25 years:

```
account = 1.7182818284590453
year = 1
while year <= 25:
    account = year*account - 1
    year = year + 1
```

Would you open the account?

I - The Chaotic Bank Society: the lesson

Let's unroll the series:

$$b_0 = e - 1$$

$$b_1 = 1b_0 - 1$$

$$b_2 = 2b_1 - 1 = 2(1b_0 - 1) - 1$$

$$b_3 = 3b_2 - 1 = 3(2(1b_0 - 1) - 1) - 1$$

You soon figure out that

$$b_n = n! \times \left(b_0 - 1 - \frac{1}{2!} - \frac{1}{3!} - \dots - \frac{1}{n!} \right) = n! \times \left(b_0 - \sum_{i=1}^n \frac{1}{i!} \right)$$


And you recall that (Maclaurin expansion)

$$e = \sum_{i=0}^{\infty} \frac{1}{i!}$$

If $b_0 = e - 1$, then b_n should converge to **zero**! You can't represent $e - 1$ exactly, error is accumulating like crazy!

II - The missing block

How many blocks of width w_b fit in a space of width w_s ? Of course $\lfloor w_s/w_b \rfloor$. For example, $\lfloor 0.65/0.1 \rfloor = 6$.



```
1 double total_width = 0.65;  
2 double block_width = 0.1;  
3 int blocks = (int) floor(total_width/block_width);
```

What happens with $w_s = 0.6$?

II - The missing block: the lesson

Computers have only **finite precision**: in base 2, the numbers $1/10$ and $6/10$ do not have a finite decimal expansion. Notice that in base 10 you have exactly the same problem with for example $1/3$.

- $\mathbb{F}(0.6) = 0.599999999999999978$.
- $\mathbb{F}(0.1) = 0.1000000000000000006$.
- In general, $\mathbb{F}(x)$ can be **slightly smaller or slightly greater** than x and this can have really nasty effects.

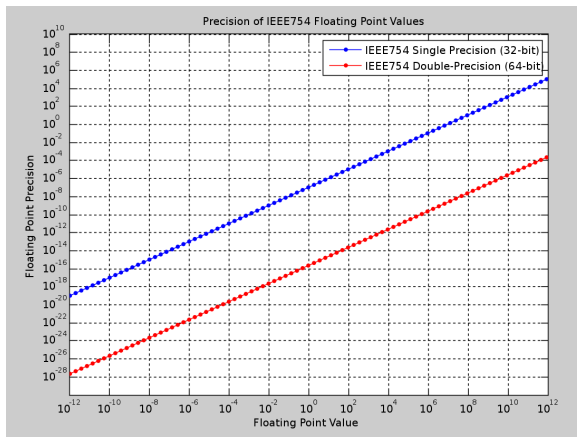
III - The odometer that stopped

Design an *odometer* that keeps track of the distance with a precision of 0.01 meters.

```
1 float meters;
2
3 void interrupt isr() {
4     if (odometer_interrupt)
5         meters += 0.01f;
6 }
7
8 int main(void) {
9     meters = 0;
10    while(1)
11        ;
12 }
```

III - The odometer that stopped: the lesson

Remember that in floating point **precision is not constant**: it depends on the magnitude of the number you are representing.



Conclusions

Knowing how a computer works allows you to

- recognize simple performance issues and not end up with super slow code
- recognize some potentially dangerous operations that will give you incorrect results

Modern engineering needs strong interdisciplinary skills, and **numerical computing is a central tool**. Even if you only write small Matlab prototype programs, computer architecture somewhat affects you.

Also if you are a theoretical person, some computer architecture is needed: you could develop the best and most beautiful algorithm in the world, but if from a practical point of view is not good no one will use it.